

95 questions, each tied to a decision

This restores the complete coded question bank from the original guide while making it easier to use in live working sessions. Nothing is ingested simply because a connector exists.

HOW TO USE THIS

Run the source map first, then the cross-cutting gates. For each system, include the owner, administrator, a daily user, and the person responsible for sensitive-data decisions. Record the answer, owner, evidence, and resulting build decision next to every code.

SECTION	COUNT	PRIMARY OUTPUT
M Source map	7	Complete source inventory and coverage disclosure
X Cross-cutting	11	Identity, vendor, legal, device, session, and offboarding decisions
D Google Drive	21	Approved folders, parsing, authority, access, and sensitivity rules
G GivingData	22	Portfolio spine, field meanings, historical mappings, and visibility
A Airtable	16	Canonical bases, reliable joins, freshness, and relationship privacy
Z Zoom Chat	12	Evidence-based decision to exclude or narrowly include chat
IM Impact model	6	Authoritative outputs, versioning, cadence, and ownership

Required artifacts: field register, source-of-truth map, access model, sensitivity tiers, Never-ingest list, first corpus, evaluation set, accepted gaps, named owners, and sign-off.

Round 0: map every place knowledge lives

Trace three past grants from first contact through renewal or close. Do not let the four named systems become an assumed boundary.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
M1	Besides Drive, GivingData, Airtable, and Zoom Chat - where else does grant information get written down? Email? Notion? Recorded calls? A spreadsheet someone maintains?	Flushes out the shadow sources before they become a credibility gap.
M2	Is there a portfolio tracker or master spreadsheet any of you keep by hand? What's in it that isn't in GivingData?	These are usually the most current view of in-flight work and the most fragile.
M3	Do you record grantee calls or internal meetings? Where do the recordings, transcripts, and summaries go, and for how long?	Zoom Cloud Recordings / transcription tools are a distinct, rich, sensitive source the brief doesn't name.
M4	Where does the impact model actually live, and who owns the numbers behind the public dashboard?	Compass must cite the authoritative impact source, and the Impact team has to be a partner, not just a stakeholder.
M5	Did the Foundation use a different grants system before GivingData? Is that data still accessible?	Determines whether "five years" is fully in GivingData or split across a migration boundary.
M6	Do you use Notion internally beyond the public Handbook? What's in it?	Likely holds strategy and meeting-notes content with a clean API - a strong early candidate to add.
M7	Are there email threads you'd expect Compass to know about - grantee decisions that only happened over email?	Sets up the (hard) question of whether and how to bring email in later.

SESSION DECISION

Rank every discovered source by value, feasibility, sensitivity, authority, historical coverage, and readiness for the first version.

Cross-cutting gates

Answer these once before connector design. Each answer changes the security or operating model.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
X1	Who is the single data owner who can approve the access model and any sensitivity-tier decision?	Every downstream "can we ingest this?" needs one accountable yes/no. Without it, discovery stalls.
X2	Is Google Workspace the identity provider for everything, or do some systems have separate logins?	Determines whether one SSO gates Compass and whether I can map permissions from one directory.
X3	What is the Foundation's written position on sending grantee and co-funder data to an AI vendor? Is there an existing agreement?	Gates the entire generation layer. If the answer is "no third parties," the model has to run in-VPC and the timeline and cost change.
X4	Which grants carry confidentiality or data-use clauses (co-funder MOUs with Ballmer Group / Annie E. Casey, individual grant agreements)?	Some grantee material may be contractually barred from secondary use. I need the list before ingestion, not after.
X5	For grantees in Colombia and Kenya, is there personal data in the reports, and has cross-border transfer ever been assessed?	Ley 1581 (Colombia) and the DPA 2019 (Kenya) have transfer conditions. Cheapest answer is usually "exclude participant PII," but I need to confirm what's in there.
X6	What does the team do today when they need an answer from grant history - walk me through the last three real instances.	The current workaround is the adoption baseline and tells me where the "obviously better" moment is.
X7	What decision, six months from now, do you want Compass to have made easier?	Anchors scope. If the answer is "renewal decisions," the corpus and the dossier view are built around that.
X8	When staff sign in to Google, does the Foundation require a second step - a code from the phone, a tap on a prompt, a physical security key? Is it required for everyone, or optional? Does a Google session stay signed in for days, or expire?	Compass signs people in with their Google account, so it gets whatever second step Google already asks for. If everyone already does a second step at Google login, Compass does not need its own separate code or authenticator app - adding one would just be an extra annoyance for no real security gain. If the second step is optional today, that's worth turning on for everyone at the Google level (one setting) rather than building it into one app. This one answer decides whether we build a second login step into Compass at all.
X9	Does the Foundation manage staff laptops and phones centrally - can IT require a screen lock, push settings, or wipe a lost device?	If yes, Compass can be set to only work on a managed device with a screen lock - a strong extra layer for a tool holding five years of sensitive grant knowledge, and it's an IT setting, not something we build. If there's no central device management, we lean on short sessions, a fingerprint/Face-ID lock on the app, and the ability to cut someone's access instantly from the server.
X10	When someone leaves the Foundation, who turns off their account, how quickly, and does that automatically cut off the connected tools too?	Compass can cut a specific person's access instantly (an admin button, plus an automatic nightly check) - but it needs to fire off the same event that disables their Google account, or a departed employee could keep a working session for a few hours. This was a real gap on the HOPE dashboard.
X11	Should Compass keep people signed in for a long time, or always sign them out after a set number of hours? Any objection to a fingerprint / Face-ID lock on the phone/installed version?	A convenience-vs-safety call for the data owner. Recommendation: sign out after ~8 hours, offer an optional fingerprint/Face-ID lock on the installed app (it's a local unlock, not a second login step), and no permanent "stay signed in". The fingerprint lock is opt-in so it doesn't get in the way of people who don't want it.

SESSION DECISION

Approve one accountable data owner, the identity and session posture, vendor terms, contractual exclusions, first decision, and revocation process.

Google Drive: orientation and structure

Drive contains narrative evidence and often the reasoning behind a formal grant record.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
D1	Is the grant-relevant content in Shared Drives or scattered across individuals' My Drive? Roughly what split?	Shared Drives have clean, inheritable permissions and survive staff departures. My Drive content is a permissions and continuity mess - if the important memos live there, the v1 corpus shrinks and I need a plan to migrate them.
D2	Which Workspace edition (Business Standard / Plus / Enterprise)? Is Google Vault in use? Are Drive Labels or DLP rules configured?	Edition determines whether I can use Labels for classification and whether admin can scope a service account. Vault means there's already a retention/e- discovery posture to align with.
D3	Will the Workspace admin provision a service account with domain-wide delegation scoped to specific Shared Drives, or do we use per-user OAuth?	Service account = clean, centrally controlled, survives staff changes. Per-user OAuth = the tool inherits each user's messy Drive and breaks when they leave. I strongly prefer the former, scoped narrowly.
D4	Can I get a folder-tree export (names, owners, item counts, last-modified) for the candidate Shared Drives before we decide scope?	I want to pick the v1 corpus from evidence, not a guess. Item counts tell me if it's 5,000 files or 500,000.
D5	Are there Shared Drives or folders that are explicitly off-limits - board, HR, legal, compensation, personnel?	These become Never-ingest or Restricted . I want them named and walled off before the first sync, not discovered later.
D6	Is there a naming convention or folder structure for grants (by org? by fund? by year? by grant ID)? How consistently is it followed?	The folder path is often the only metadata a Drive doc has. If "Care Economy/2025/Carina/..." is reliable, I get fund, year, and org for free. If it's chaos, I lean harder on entity resolution from content.
D7	Do documents follow templates (a standard proposal template, a standard report template)? Where are the master templates?	Templates give me predictable structure to extract from - "Projected lifetime earnings" is always in section 3. Free- form docs need more work and are less reliable.
D8	For Google Sheets used as impact models - is there a standard layout, and which cells hold the North Star ratio, projected earnings delta, and participant count?	Spreadsheets are the hardest to ingest well. I need to know if there's a consistent structure I can target or if every model is bespoke.
D9	Are grant reports mostly native Google Docs, uploaded PDFs/Word, or scanned PDFs? Roughly what mix, and does it change for older grants?	Native Docs extract cleanly. Scanned PDFs need OCR and a confidence check. If five-year-old reports are all scans, backfill quality drops and I budget for it.
D10	Do you rely on Google Docs revision history / comments as a record of decisions or feedback?	If yes, I may need to capture comment threads, not just final text - that's where "we pushed back on their target" often lives.

SESSION DECISION

Approve the candidate Shared Drives, access method, file and template scope, comments policy, and parsing requirements.

Google Drive: history and sensitivity

Test the older, messier, and more sensitive material before treating five-year coverage as reliable.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
D11	Has the proposal or report template changed over the last five years? When, and how different is old vs. new?	A field that moved or got renamed breaks naive extraction. I need to map old/new so a 2021 report and a 2026 report answer the same question.
D12	Was there ever a migration into Drive (from Box, Dropbox, a previous shared drive, a departed employee's account)? Any known gaps or duplication from it?	Migrations leave duplicates, broken links, and missing metadata. If one happened, I plan dedupe and reconciliation around it.
D13	Is it common to have multiple versions of the same report in the same folder ("v2", "FINAL", "FINAL_clean")? Which one is real?	This is the "no single source of truth" problem in miniature. I need the team's rule for which copy counts so the dedupe layer picks the authoritative one.
D14	Are there grantee reports in Spanish (Colombia) or other languages? Are they translated, and if so is the translation in the same folder?	Determines multilingual retrieval scope and whether I index the original, the translation, or both (linked).
D14a	Separate from the documents: does any staff member work primarily in a non-English language and expect to ask questions and read answers in that language (not just have non-English documents rendered into English for them)?	Decides whether Compass needs a per-user language preference and answer-in-the-reader's- language rendering. Today translation is one- directional (any language to English at intake). If no staff gap, the current design is correct; if yes, it's a scoped addition, not a redesign.
D15	Roughly what fraction of grants have a complete document set in Drive (proposal + each required report)? Where are the known holes?	Sets expectations for the gap report and tells the team where Compass will honestly say "we don't have that."
D16	Do grant reports contain named information about program participants (people served by grantees - names, stories, contact info, immigration or justice-system status)?	This is the highest-sensitivity content in the whole system. Default plan: detect and redact at intake, never index participant identifiers in v1. I need to confirm how prevalent it is.
D17	Are there documents with candid internal assessments of a grantee's leadership or viability that would damage the relationship if surfaced casually?	These are Restricted or Programs-only . A synthesis answer must never quote "the ED is difficult to work with" back to a wider audience.
D18	Do any documents contain grantee financials, audits, or bank details provided in confidence?	Restricted . Financial-account data gets the same treatment as PII.
D19	Within a Shared Drive, is permission uniform, or are there restricted subfolders with tighter sharing?	The connector must respect folder-level ACLs, not just drive-level. If the model is "drive = team, but /board is locked," the permission map has to capture that.
D20	Does the Foundation have retention rules for grant documents (keep X years after close)? Any active legal holds?	Compass references records; it must not become a way to resurrect something that was properly deleted, or to circumvent a hold.

SESSION DECISION

Approve template and language mappings, authoritative-version rules, sensitive-content handling, folder-level access behavior, and retention alignment.

GivingData: orientation and structure

Treat GivingData as the likely portfolio spine, but verify what its custom fields actually mean and how complete they are.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
G1	Does the Foundation's GivingData plan include API access? What endpoints, what auth (API key / OAuth), what rate limits?	Determines whether I build a real incremental connector or fall back to scheduled report + document exports. Changes freshness and effort materially.
G2	If no API: what scheduled exports are possible - full CSV of grants/orgs/payments/requirements, plus a bulk document export from the grantee portal?	The fallback has to be complete and re-runnable. A partial export creates silent gaps.
G3	Who is the internal owner (likely the Grants Manager)? Who administers custom fields and workflow?	I need one person who knows why every custom field exists and which ones are abandoned.
G4	How many grants / requests are in the system total, and from what date? Is pre-Foundation-era or pre-migration data in there?	Confirms corpus size and where "five years" actually starts.
G5	Do all Programs staff currently see all grant records, or is visibility restricted by role / portfolio / fund?	Compass must mirror this. If a program officer can't see another portfolio's grants in GivingData, they can't see them through Compass either.
G6	Walk me through the grant record: standard fields (amount, dates, status, type) plus every custom field. Which are required, which are optional, which are effectively dead?	The custom fields are where the Foundation- specific meaning lives - thesis area, North Star projection, cohort. I need to know which are trustworthy and which are 40% filled.
G7	How is projected impact captured - a North Star ratio field? projected annual/lifetime earnings delta? projected participants reached? Are these entered at approval and never revisited, or updated?	This is half of every "how did they do vs. plan" answer. I need to know it's a real field and understand its vintage.
G8	How is reported / actual impact captured - structured fields on the requirement record, or only inside an uploaded report document?	If actuals are structured, cross-system answers are clean joins. If they're buried in PDFs, the pipeline has to extract them and the entity graph does more work.
G9	How do Requirements (reports, check-ins) work - schedule, due dates, status, submitted documents, internal review notes? Is there a "report received / reviewed / approved" state?	Requirements drive the completeness reconciliation: "the schedule says a Q3 report was due; is it here?"
G10	What's in the review / scorecard data - the impact-scorecard rubric dimensions, scores, reviewer notes? Is it in structured fields or attached memos?	The Foundation's diligence rationale. Valuable for "why did we fund / decline this" - but often sensitive (Programs-only).
G11	Organizations and Contacts: what's the schema, and how does it overlap with what's in Airtable? Which system is authoritative for org profile and relationship data?	Directly informs the source-of-truth map. If both hold org data, I need the rule for which wins.
G12	Payments / financials: payment schedule, amounts, dates, GL/budget codes, fund allocation. How much financial detail should Compass index vs. leave to Finance's systems?	Grant-dollar data is in scope, but I don't want Compass reproducing the general ledger. CFO/Controller should set the line.
G13	How are Funds / Portfolios / Cohorts modeled (AI for Economic Opportunity, Powering Economic Opportunity, Learning for Action)? Can every grant be tied to one?	The cohort is a primary filter and a common query axis ("pull every outcome for the AI Fund cohort").

SESSION DECISION

Approve the connector path, stable identifiers, core fields, impact semantics, requirements model, financial boundary, and portfolio taxonomy.

GivingData: history and sensitivity

Historical meaning, migration quality, duplicate records, declined applicants, and review notes require explicit rules.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
G14	Was there a migration into GivingData (from Fluxx, Foundant, spreadsheets, a prior system)? What came over clean and what didn't?	Migrated grants often have blank custom fields, lost attachments, or mangled dates. The reconciliation layer needs to know.
G15	Have field definitions or the workflow changed over five years - a renamed status, a retired custom field, a changed rubric?	Same issue as Drive template drift. A status that meant one thing in 2022 and another in 2026 corrupts filtering unless mapped.
G16	Are there duplicate organization records (same grantee entered twice, or as "Org" and "Org, Inc.")? Duplicate grants?	Duplicates in the spine are the worst case - they double-count in portfolio answers. Dedupe rules start here.
G17	How complete is thesis-area / geography tagging across the full grant history?	If 30% of grants are untagged, "what have we funded in credential completion" quietly under- reports. The gap report must surface this.
G18	For grantee-portal uploads - do grantees sometimes upload the wrong file, a draft, or a duplicate of something also in Drive?	Cross-system dedupe: the portal PDF and the Drive Doc are often the same report. One is authoritative.
G19	Do grantee-portal submissions come with confidentiality terms? Is there language in the portal or grant agreement about how the Foundation uses submitted materials?	If the terms don't cover internal AI-assisted analysis, the privacy notice / agreement needs updating before ingestion (§6.8).
G20	Which fields or record types are sensitive - reviewer scorecards, internal notes, declined-applicant records, anything about organizational instability?	Sets tiers. Declined-applicant diligence is usually Programs-only and the applicant name may need redaction in a general view.
G21	Are there co-funded grants where another funder's data or terms are embedded in the record?	Ties to X4 - co-funder confidentiality. Flag these records.
G22	Does GivingData hold participant-level data anywhere (some grantees report line-level outcomes)?	Same as D16 - highest sensitivity. Confirm and plan redaction.

SESSION DECISION

Approve historical mappings, deduplication, tagging coverage, portal-document authority, confidentiality terms, and Programs-only fields.

Airtable: ownership, access, and structure

Separate canonical relationship data from mirrors, personal experiments, formulas, and overlapping bases.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
A1	How many bases are relevant, and who owns each? Is there one "Relationships" base or several overlapping ones?	Airtable sprawl is common. I need the canonical base(s) and to know which are personal experiments vs. team-of-record.
A2	Can I get a read-only API token (or OAuth) scoped to those bases? Note Airtable's ~5 requests/sec limit and that attachment URLs expire.	Determines connector design - I have to pull attachments promptly and cache them, not store expiring links.
A3	Who are the collaborators on each base, and at what permission level (owner / editor / commenter / read)?	Feeds the permission map. If a base is shared org-wide as read-only, its content is Team ; if it's shared with three people, it's tighter.
A4	Are there automations, sync-ins, or integrations on these bases (Airtable syncing from GivingData, forms feeding records, Zapier)?	Tells me if Airtable is a source or a mirror. A field synced from GivingData shouldn't be treated as independent data.
A5	Walk me through each table: Organizations, Contacts, Pipeline/Deals, Interactions, plus any others. What are the fields and the linked-record relationships between tables?	Linked records are how Airtable models the graph. I want to reuse that structure in entity resolution rather than rebuild it.
A6	Is there a field that holds the GivingData grant ID or org ID? Any reliable cross-reference?	If yes, entity resolution between Airtable and the spine is trivial. If no, I'm fuzzy-matching on org name and it's error-prone.
A7	How is the interaction / touchpoint log structured - date, type, attendees, notes, linked org? Is it kept up reliably or sporadically?	Interaction history powers the dossier's "last contact / relationship state." If it's sporadic, I present it as partial, not authoritative.
A8	What tags / single-selects / multi-selects exist for thesis area, geography, stage, source, priority? Do they match GivingData's vocabulary?	Controlled-vocabulary mapping. If Airtable says "workforce" and GivingData says "Future of Work," the normalization layer reconciles them.
A9	Which fields are formulas, rollups, or lookups (derived) vs. directly entered?	I ingest entered data as fact; derived fields I either recompute or mark as derived so they're not treated as independent evidence.

SESSION DECISION

Approve the canonical bases, read-only token scope, table graph, GivingData joins, controlled vocabularies, and derived-field handling.

Airtable: quality and sensitivity

Relationship information becomes misleading when stale and harmful when private context reaches the wrong audience.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
A10	Are there duplicate org or contact records? How does the team currently handle merging?	Airtable has no strong dedupe. Duplicates here create split relationship histories.
A11	How current is the data - when was the pipeline / interaction log last meaningfully updated? Are there stale records for people who've left grantee orgs?	Stale CRM data surfaced confidently is worse than no data. The dossier should date every relationship fact.
A12	Has the base structure changed - tables renamed, fields added/removed, a big restructure? Any prior base that was abandoned?	Abandoned bases sometimes hold history the current one lost. Worth checking before scoping.
A13	Do people keep private notes in Airtable fields (personal impressions, "don't trust their numbers")?	Same as D17 - candid content that needs a tier and careful handling.
A14	Which fields hold contact PII - personal email, personal phone, home address, LinkedIn, personal notes?	Contact PII is redacted unless an answer genuinely needs it. A general "who's our contact at X" can return a role and work email, not a cell number.
A15	Are there funders / donors / prospects in here whose information is confidential (a donor who wishes to stay anonymous)?	Donor confidentiality is serious for a foundation. Those records may be Restricted or Never- ingest .
A16	Any view-level or field-level restrictions currently in place that express "only these people should see this"?	Mirror them. Airtable's own permission model is the team's stated intent.

SESSION DECISION

Approve freshness standards, duplicate handling, historical base coverage, contact-field boundaries, donor confidentiality, and restricted views.

Zoom Chat: prove that inclusion is justified

Zoom Chat is the highest-risk named source. These questions support a decision to defer it unless unique value, retention, employee expectations, and counsel approval align.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
Z1	Is this Zoom Team Chat (persistent channels) or just in- meeting chat? What's the message-retention setting at the account level - 30 days, 1 year, indefinite?	This is the first gate. If retention is 1 year, "five years of team conversations" does not exist and the whole premise for this source is wrong. Likely outcome.
Z2	Does the Foundation have admin rights to create a Zoom Server-to-Server OAuth app with chat_message:read:admin scope? Or is chat history only reachable per-user?	Admin scope is the only practical way to ingest at team scale. Without it, this source is effectively closed.
Z3	Is there a compliance/archiving integration already (Global Relay, Smarsh, Zoom's own archiving)?	If archiving exists, there's already a governed copy of chat with a retention and e-discovery posture - ingest from that, not the live API, and inherit its governance.
Z4	How many channels, and what's the split of public vs. private vs. DMs? Which channels are actually about grant decisions vs. social / logistics?	Scopes the useful surface. Probably a handful of program channels matter and everything else is noise or off-limits.
Z5	In the grant-relevant public channels, how are decisions actually recorded - a clear "we've decided X" message, or a scattered thread you have to read in full?	Determines whether thread-level chunking can capture a decision with its context, or whether the signal is too diffuse to be useful.
Z6	Do important decisions in chat also get written up somewhere in Drive or GivingData?	If yes (usually the answer), the marginal value of ingesting chat is low and the privacy cost isn't worth it - which supports the deferral.
Z7	Are files and links shared in channels that aren't stored anywhere else?	Occasionally chat is the only home for a useful artifact. Worth knowing, but usually recoverable another way.
Z8	What have staff been told about their Zoom Chat - is it understood to be searchable / archived, or do people treat it as ephemeral and private?	The single most important question. If people speak candidly believing it's ephemeral, ingesting it is a trust violation regardless of the legal position.
Z9	Would ingesting public-channel history require notifying staff and offering an opt-out? Is there a works-council / employment-law consideration?	I would not ingest without explicit disclosure. This question sets the process.
Z10	Are DMs and private channels ever on the table?	My answer is no - permanently. This question is to confirm leadership agrees and put it in writing.
Z11	Does chat contain personnel, compensation, or performance discussion, or candid personal opinions about grantees and partners?	Almost certainly yes. This content is exactly what must never surface in a team synthesis tool.
Z12	What are the e-discovery / legal-hold implications of creating a second indexed copy of chat?	A searchable index of chat is discoverable. Counsel has to weigh in before this source is touched.

SESSION DECISION

Default to exclusion. If a narrow case proceeds, use approved public channels only, explicit disclosure, an existing governed archive where possible, and permanent exclusion of direct messages and private channels.

Impact model: protect the meaning of the numbers

Compass should cite authoritative impact outputs and their versions. It should not quietly reimplement the Foundation's methodology.

#	ACTUAL DISCOVERY QUESTION	WHY I'M ASKING / DECISION IT CHANGES
IM1	Where does the model physically live - a Tableau data source? a set of Google Sheets? a dbt / Python pipeline into a warehouse (BigQuery)? a notebook with no saved output?	Decides whether there's anything to connect to. Sheets / BigQuery to a clean read connector. Tableau to connect to the data source under the dashboard, not Tableau's API (that's for rendering dashboards, not clean data). A notebook with no persisted output to Compass can only ever cite the memo, and we say so in the coverage line.
IM2	What's the grain - one North Star number per grant? per cohort? does it change with each model version, and when is it re-run?	Compass has to cite the right number. If numbers are versioned, every impact citation must carry the model version + date, and "projected 94x under v3, restated to 96x under v4" becomes a conflict-surfacing case (shown, both values, not silently merged) rather than a contradiction the reader has to catch.
IM3	Projected vs. reported - does the model hold both the projection and the actuals, or is "actual / reported" only ever inside a grantee's report?	The single most important distinction in the product. If actuals live only in Drive reports, Compass links the model's projection to the report's actuals itself and shows the delta (this is exactly what the Grantee dossier's "projected vs. reported" row does today). If the model holds both, Compass cites the model for both and the join is trivial.
IM4	Who owns the numbers, and who resolves it when Compass's answer and someone's dashboard disagree?	Someone will ask "why did Compass say 94x when my dashboard says 96x." A named owner means model-vs- memo conflicts have a person who adjudicates. No owner means every impact answer is contested.
IM5	Is the model itself sensitive - are the cost assumptions, discount rates, or unpublished inputs confidential?	Sets the tier. The model's outputs might be team-tier while an assumptions tab is Restricted to the connector ingests the outputs and stubs the assumptions.
IM6	Refresh cadence - nightly extract, quarterly recompute, annual?	Freshness is one of the four inputs to Compass's confidence score. A quarterly model means a 4-month-old impact citation is normal, and the freshness scoring has to know the model's natural cadence - the same way per-grant reporting cadence is handled.

SESSION DECISION

Approve the physical source, grain, joins, projected-versus-reported meaning, model version, owner, refresh cadence, and confidential assumptions.

FINAL SYNTHESIS

For each fact type, record the authoritative source, primary join key, conflict owner, freshness expectation, sensitivity tier, and accepted coverage gap. Leadership then approves the first corpus, Never-ingest list, evaluation set, and release conditions.

Operating lesson: "We do not have that" is a feature. People trust the system more when it reliably admits gaps than when it produces a polished but unsupported answer.