

The interview in one page

Your advantage is not that you built a chatbot. You understood the operating problem, built a contained test, and can explain how to turn fragmented experiments into shared capability.

THE CORE THESIS

GitLab Foundation needs someone who can connect programs, operations, impact measurement, and technology without weakening human judgment or trust.

Compass demonstrates that combination. It connects evidence across systems, checks access before retrieval, cites the source, surfaces disagreement, and keeps grant decisions with people.

Your opening

"When the role was described to me, I heard a need for more than another AI tool. I heard a need to bring the team's experiments and information into a shared, trustworthy system that program officers, operations, and impact staff can all use. That is the problem I scoped Compass against. The case study is the answer to the assignment, and I used a fictional demo to pressure-test the assumptions behind it."

WHAT THEY NEED

- **Diagnosis before automation:** observe the work and prioritize with Elicia.
- **Consolidation:** make existing experiments shared and supportable.
- **Maintainability:** documentation, tests, ownership, and clear handoff.
- **Responsible AI:** access, provenance, human review, and honest refusal.
- **Delivery:** tangible wins during the fellowship, with a path to permanence.

WHAT YOU BRING

- **The ask under the ask:** you start with the decision, not the feature.
- **Builder plus translator:** architecture and plain-language adoption.
- **Foundation experience:** Dash proves you know this operating shape.
- **Safety as behavior:** permissions and refusal are built into the workflow.
- **Visible progress:** working software, tests, documentation, and iteration.

The sentence to remember

"I find the decision underneath the request, translate it into a system people can understand, and build it so the team can trust, use, and maintain it."

Why this role matters now

The Foundation is scaling its capital, team, partnerships, and AI activity. Fragmented information becomes a larger operational risk at each step.

CHANGE	WHAT IT CREATES	YOUR RESPONSE
Grantmaking moving toward larger, fund-based and multi-funder structures	More grants, reporting obligations, data-use constraints, and cross-fund questions	A permissioned evidence layer that preserves each source's meaning and owner
AI experiments emerging across the organization	Useful local tools, uneven skill levels, limited visibility, and maintenance risk	Build on what exists and create shared patterns, documentation, and governance
Grantmaking-as-a-service ambition	Internal capabilities may later need to work for peer funders	Modular connectors and policies, but only after the internal workflow proves useful
Fast organizational growth	Institutional knowledge cannot depend on the longest-tenured person	Cited knowledge access, onboarding support, and visible decision history

What success means in their words

Diagnose first. Shadow teams, inventory tools, map repeated workflows, and bring Elicia a prioritized shortlist.

Deliver visible wins. Ship a small useful slice, show progress, gather corrections, and widen only when the evidence supports it.

Bring the whole team. Support power users and reluctant users through examples, playbooks, training, and patient coaching.

Leave it stronger. Tests, runbooks, deploy guidance, ownership, access rules, and decisions that do not live only in your head.

YOUR 90-DAY ANSWER

Listen and map. Learn the real questions, owners, systems, and risks. **Prioritize together.** Rank opportunities by usefulness, effort, sensitivity, and adoption. **Ship and learn.** Deliver a narrow workflow, measure whether it helps, document it, and decide what deserves expansion.

Operating principle: Move quickly by making the next decision clear. Do not skip the decision that protects people, data, or measurement integrity.

Ellie Bertani and Elicia Wilson

Ellie is testing ambition and judgment. Elicia is testing whether she can trust you with an ambiguous, cross-team operating problem.

Ellie Bertani

President and CEO

Remit. Strategy, organizational ambition, economic-mobility impact, and the Foundation's results orientation.

What they are testing. Can you raise the organization's ambition while treating AI as a means, not the mission? Can you move with both confidence and humility?

Lead with. Compass as infrastructure for moving faster with integrity as the portfolio and partnerships scale. Show honest failure: cited answers, visible conflict, and refusal when evidence is weak.

Avoid. Model hype, a long product roadmap, loose mission language, or being overly modest about what you can deliver.

Likely questions. What would you do in the first two weeks? How do you know a tool is trustworthy? Where could AI most change what the Foundation can do? What failure could set the organization back?

Elicia Wilson

Chief Operating Officer and hiring manager

Remit. Operations across grantmaking, impact advisory, capital mobilization, the Handbook, and the remote operating model.

What they are testing. Can she hand you a messy problem and receive a finished, documented, maintainable result without micromanaging you?

Lead with. Your discovery method, prioritization rhythm, handoff artifacts, security controls, and visible progress. Use Dash as proof that you have scoped a foundation workflow before.

Avoid. Treating documentation or security as polish, criticizing existing experiments, or sounding as if you need detailed instructions before moving.

Likely questions. What would you do in the first 30 days? How does your work survive your departure? How do you decide what not to automate? What do you need from each team?

HOW TO ADJUST YOUR ENERGY

With Ellie: be ambitious, direct, and first-principles. Show what becomes possible at scale. **With Elicia:** be concrete, calm, and operational. Show how you prioritize, report progress, document decisions, and reduce dependency on one person.

Strong proof to use

- **Dash:** a live foundation workflow covering more than \$7M, 41 grants, and 20 cohorts per year.
- **Compass:** the requested scope plus a fictional-data demo used to test permissions, citations, conflicts, and refusal.
- **Your scope-pivot story:** you protected the core scoring system and moved the broader request into a responsible future phase.
- **Your documentation habit:** tests, runbooks, deployment guidance, security controls, and a clear promotion gate.

Best handoff line: "I do not want the value of a tool to depend on me remaining the only person who understands it."

IF ELICIA ASKS: FIRST 30 DAYS?

"I would listen across teams, map the repeated workflows and current AI experiments, identify the owners and access boundaries, then bring you a prioritized shortlist. We would choose the first workflow together based on usefulness, risk, effort, and whether we can measure the improvement."

IF ELLIE ASKS: WHERE CAN AI CHANGE THE ORG?

"AI can shorten the distance between what the Foundation already knows and the next responsible action. The opportunity is faster access to evidence, stronger organizational memory, and more consistent learning across a larger portfolio."

Matt Zieger and Tamsin Chen

Matt is testing usefulness for grantmaking. Tamsin is testing whether the system makes impact claims more rigorous or less rigorous.

Matt Zieger

Chief Programs and Partnerships Officer

Remit. Grantmaking, partnerships, the AI for Economic Opportunity portfolio, program-officer workflows, and emerging grantmaking services.

What they are testing. Does this make grantmaking sharper and faster? Do you understand grantees and program work, not only the technology?

Lead with. The renewal-preparation view: grant cycle, projected versus reported results, relationship history, source links, gaps, and who to ask. Explain how it reduces repeated research without replacing judgment.

Avoid. Implying that Compass replaces GivingData or existing experiments. Do not minimize participant data, candid assessments, or declined-applicant sensitivity.

Likely questions. How does this change a program officer's day? What would you build first? How do you protect candid notes? How would you measure whether it helps?

Tamsin Chen

Director of Impact

Remit. Impact measurement, the North Star model, lifetime-earnings estimates, methodology, and impact advisory.

What they are testing. Will AI preserve provenance, attribution, counterfactual reasoning, comparability, and model versioning, or will it make weak evidence sound strong?

Lead with. Compass retrieves and cites impact information. It does not invent or recompute an impact figure. It labels dates, methods, model versions, coverage, and disagreement.

Avoid. Saying the AI assesses impact, scores grants, or figured out a number. Do not blur outputs, outcomes, and impact.

Likely questions. How do you stop a weak number from looking strong? What happens when sources disagree? How would you handle the impact model? How can someone verify a number?

THE PROGRAMS AND IMPACT BRIDGE

A useful system helps a program officer find the evidence. A trustworthy system also preserves what the evidence means. Your design connects those needs: structured records for dates and amounts, narrative records for context, direct citations for verification, and a human owner for interpretation.

Language that will land

- "The system can phrase the evidence, but it does not make the grant decision."
- "Confidence describes coverage, freshness, source agreement, and citation completeness. It is not the model saying it feels confident."
- "If the memo and the official record disagree, the answer shows both, identifies their dates, and names who should resolve it."
- "Compass would cite the economic-impact model and its version. I would not reimplement Tamsin's methodology inside a chatbot."

Why you are unusually well suited

Your differentiator is the combination. You understand nonprofit operations, can build the technical system, can explain it to mixed-skill users, and naturally design for trust.

YOUR INSTINCT	WHY IT MATTERS HERE	PROOF
Find the ask under the ask	The role requires workflow discovery before choosing an AI intervention.	Dash discovery and the partner-scouting scope decision
Translate across functions	Programs, operations, impact, leadership, and technical contributors need one shared language.	You pair architecture with plain-language explanation and user flows.
Build to learn	The Foundation values learning by doing, but needs experiments that create evidence.	Compass fictional-data prototype, QA findings, and iterative hardening
Design the failure path	Sensitive data and impact claims require refusal, conflict handling, and provenance.	Permission checks, source links, no-answer behavior, human review
Think about adoption	The team spans power users and people new to AI.	Role and cycle context, playbooks, training, and coaching instincts
Finish and hand off	Their feared failure is non-delivery or a tool only its creator can operate.	Live builds, tests, runbooks, deployment guidance, and promotion criteria

FULL POSITIONING

"I have worked on both sides of this problem. I understand the human and operational reality of nonprofit work, and I can translate that reality into a secure, maintainable technical system. I begin with the decision, the people affected, and the evidence they need. Then I build enough to test the idea, document it so others can carry it forward, and explain it in language the whole team can use."

SHORT POSITIONING

"I am a translator between programs, operations, leadership, and technology. I find the ask under the ask, turn it into something practical, and make sure the result is safe, understandable, and usable after I leave."

How to describe your intensity: "I do not abandon a project at the prototype stage. I keep testing, documenting, and strengthening it until the team has something it can actually use. The team, not my enthusiasm, determines what moves into production."

Your proof bank

STORY	WHAT IT PROVES	USE IT FOR
Dash	Foundation workflow knowledge, role-based access, grant-cycle operations, and a working AI query layer	Discovery, adoption, permissions, and handoff
Flood Voice	You can move from an urgent public problem to a working solution and external interest	Speed, ambiguity, initiative, and public impact
RhythmWell	You can build thoughtfully in a high-stakes environment and communicate with institutional partners	Trust, health context, stakeholder communication
Scope pivot	You can protect the useful core while acknowledging a larger opportunity	Prioritization, feedback, restraint, and collaboration

Questions, landmines, and close

Ask questions that expose priorities and operating reality. Close by connecting your delivery record to the exact problem they described.

ASK ELICIA

- Which workflow is creating the most operational drag right now?
- As the fund model grows, what do you expect to break first?
- What would visible progress from this fellow look like to you?

ASK ELLIE

- Two years from now, what should the Foundation be able to do that it cannot do today?
- Where is philanthropy still moving too slowly?

ASK MATT

- Which recurring program decision most needs better information?
- What should remain distinctly human as grantmaking becomes more AI-supported?

ASK TAMSIN

- What measurement mistake would you most want designed out of an internal tool?
- How should a knowledge system cite the evolving impact model?

Landmines to avoid

Do not imply certainty about internal systems. Say what is based on public research or fictional data and what you would confirm in discovery.

Do not frame Compass as replacing the team's work. Credit existing experiments and describe the shared layer you would build with their creators.

Do not oversell AI. Talk about access, speed, evidence, and better workflows. Keep the human on the decision.

Do not use loose impact language. Separate projected, reported, and modeled values. Preserve the source, date, and methodology.

Do not lead with a giant roadmap. Lead with the smallest valuable workflow and the evidence required to widen it.

Do not make security a final slide. Explain how permissions and data boundaries change the architecture from the beginning.

CLOSING

"What excites me about this role is that the Foundation already has the ambition, experiments, and domain expertise. I can help turn that energy into shared systems that people can trust and maintain. I have built this general problem shape before, I know how to learn from the people closest to the work, and I know how to move from a useful test to something the organization can carry forward."

Reference points from the original briefing

GitLab Foundation website and public Handbook; Applied AI Fellow posting; public profiles and interviews for Ellie Bertani, Elicia Wilson, Matt Zieger, and Tamsin Chen; the candidate's recruiter and first-round notes; Dash, RhythmWell, Flood Voice, and Compass project materials.